

Mathematical Formulation of the Dispatch Optimization Algorithm

bess-stack-de

Degradation-aware rolling-horizon dispatch for a grid-connected BESS in the German market

V1.0

1 Scope and Conventions

This document states the optimization model implemented in `./src/bess_stack/model/dispatch.py`. The project models

a degradation-aware, rolling-horizon dispatch optimizer for a single grid-connected BESS in the German market, co-optimizing day-ahead arbitrage, FCR capacity, and aFRR capacity + energy, that outputs an annual revenue stack and feeds a project-level IRR calculation.

Section 2 describes what is implemented. Section 3 describes what is not. Every equation in Section 2 corresponds to a statement in `solve_window`; no equation there describes intended behaviour. Section 3 formulates the reserve products from the parameters already present in `scenarios/reference.yaml`, none of which the code reads today. The distinction is maintained deliberately: a formulation that does not say which half is running is not a specification but an advertisement.

1.1 Units and sign conventions

Power is in MW, energy in MWh, prices and costs in EUR/MWh, and capacity prices in EUR/MW/h. One time step spans Δt hours. Charging and discharging power are separate non-negative variables rather than one signed variable, which is what allows the complementarity condition of constraint group B1 to be written as a linear constraint. State of charge is carried as an energy E_t in MWh rather than a percentage; the scenario file states its bounds as fractions and the model multiplies them by the nominal capacity on load.

2 Day-Ahead Arbitrage Model (Implemented)

2.1 Index sets

Symbol	Description
$t \in \mathcal{T}$	Time steps within one optimization window, $\mathcal{T} = \{0, 1, \dots, N - 1\}$
$\mathcal{B}_j$	Steps belonging to day-ahead product block j , so that $\bigcup_j \mathcal{B}_j = \mathcal{T}$

2.2 Parameters

Each parameter names the field it is read from, so that the document and the configuration cannot drift apart unnoticed. Source fields are dotted paths into `scenarios/reference.yaml`; those beginning with a stream name (`day_ahead`, `fcr`, `afrr`) are relative to `revenue_streams`.

Symbol	Description	Unit	Source field
Δt	Length of one time step	h	<code>market.resolution_minutes</code>
π_t	Day-ahead price in step t	EUR/MWh	price series argument
κ	Spread retention, $1 - \text{haircut}$	–	<code>day_ahead.spread_haircut</code>
f	Trading fee per MWh traded	EUR/MWh	<code>day_ahead.fees_eur_per_mwh</code>
g^c	Network fee on energy charged	EUR/MWh	<code>grid.charging_fees_eur_per_mwh</code>
g^d	Network fee on energy discharged	EUR/MWh	<code>grid.discharging_fees_eur_per_mwh</code>
c^{deg}	Marginal degradation cost	EUR/MWh	<code>degradation.marginal_cost_eur_per_mwh</code>
$\bar{P}$	Binding power cap	MW	<code>min(battery.power_mw, grid.connection_limit_mw)</code>
E^{nom}	Nominal usable energy	MWh	<code>battery.energy_mwh</code>
η^c	One-way charging efficiency	–	<code>battery.eta.charge</code>
η^d	One-way discharging efficiency	–	<code>battery.eta.discharge</code>
σ	Self-discharge, fraction of E^{nom} per day	1/d	<code>battery.eta.self_discharge_per_day</code>
P^{aux}	Parasitic auxiliary load	MW	<code>battery.aux_load_mw</code>
L	Idle energy loss per step, see (2.5)	MWh	derived
E^{min}	Lower energy bound	MWh	<code>battery.soc.min</code> $\times E^{\text{nom}}$
E^{max}	Upper energy bound	MWh	<code>battery.soc.max</code> $\times E^{\text{nom}}$
E_{init}	Energy at window start	MWh	carried between windows
E^{term}	Required energy at window end	MWh	<code>dispatch.terminal_soc_fraction</code> $\times E^{\text{nom}}$
m	Time steps per day-ahead block	–	<code>day_ahead.product_resolution_minutes</code>
H	Optimization horizon	h	<code>dispatch.horizon_hours</code>
C	Committed step of the rolling horizon	h	<code>dispatch.commit_step_hours</code>

2.3 Decision variables

Variable	Domain	Unit	Description
P_t^c	$[0, \bar{P}]$	MW	Charging power in step t
P_t^d	$[0, \bar{P}]$	MW	Discharging power in step t
u_t	$\{0, 1\}$	–	Mode indicator, 1 = charging permitted, 0 = discharging permitted
E_t	$[E^{\text{min}}, E^{\text{max}}]$	MWh	Stored energy at the end of step t

2.4 Objective function

The model maximizes net trading revenue over the window:

$$\max \sum_{t \in \mathcal{T}} \left[(\pi_t \kappa - f - g^d - c^{\text{deg}}) P_t^d \Delta t - (\pi_t \kappa + f + g^c) P_t^c \Delta t \right] \quad (2.1)$$

Equation (2.1) is the sum over all steps of revenue from selling less the cost of buying. Its terms are:

- $\pi_t \kappa$ — the price actually captured. The haircut κ shaves the traded spread to represent the gap between the price used in the model and the price a real bid achieves.
- f, g^c, g^d — per-MWh fees. Note the asymmetry of sign: fees subtract from the sale price and add to the purchase price, so they narrow the spread from both ends rather than acting as a levy on net position.

- c^{deg} — the marginal cost of degradation, charged per MWh discharged. This is not a cash cost and nobody invoices it. It is the shadow price of consuming battery life, and it is what makes the model degradation-aware rather than merely degradation-reporting: a cycle whose spread does not cover c^{deg} is declined. Its admissible range is bounded above by the cost of making good the capacity one cycle consumes,

$$c^{\text{deg}} \leq \phi \cdot \kappa^{\text{aug}} \cdot 1000,$$

where ϕ is `degradation.cyclic_fade_per_full_cycle` and κ^{aug} is `degradation.augmentation.cost_eu`. Nominal energy cancels out of the ratio, because fade per cycle is a fraction of beginning-of-life capacity and one equivalent full cycle discharges exactly that. `config.py` enforces the bound. There is no corresponding lower bound: establishing one needs a discount factor keyed to the augmentation date, which depends on the cycle count, which is an outcome of dispatch rather than an input to it. Satisfying the bound therefore shows the value is not indefensibly high, not that it is right.

- Δt — converts power in MW to energy in MWh over the step.

Because degradation is charged on discharge only, throughput is measured as discharged energy throughout this document.

2.5 Operational constraints

B1 — Charge/discharge complementarity and power limits

A battery may not charge and discharge in the same step. The restriction is enforced through the binary mode indicator, which simultaneously imposes the power limit:

$$P_t^c \leq \bar{P} u_t \quad \forall t \in \mathcal{T} \quad (2.2)$$

$$P_t^d \leq \bar{P} (1 - u_t) \quad \forall t \in \mathcal{T} \quad (2.3)$$

where $\bar{P}$ is the smaller of the battery power rating and the grid connection limit, and u_t selects the mode.

This binary is not a modelling nicety and must not be relaxed to obtain a linear program. Under the relaxation, at sufficiently negative prices the optimizer charges and discharges simultaneously: the two positions cancel in the objective at zero net cost, while round-trip losses destroy stored energy. Destroying energy for free is valuable when prices are negative, because it lets a full battery keep buying. The behaviour is an artefact of the relaxation rather than a dispatch strategy, and the binary is what forbids it.

B2 — State-of-charge dynamics

Stored energy evolves under charging, discharging and idle losses:

$$E_t = E_{t-1} + \eta^c P_t^c \Delta t - \frac{P_t^d \Delta t}{\eta^d} - L \quad \forall t \in \mathcal{T} \quad (2.4)$$

with $E_{-1} = E_{\text{init}}$, and the idle loss per step

$$L = \left(\frac{\sigma E^{\text{nom}}}{24} + P^{\text{aux}} \right) \Delta t. \quad (2.5)$$

where η^c multiplies the charging term and η^d divides the discharging term. The asymmetry is deliberate and is the correct treatment of one-way efficiencies: of the energy drawn from the grid only a fraction η^c is stored, whereas delivering P_t^d to the grid requires drawing P_t^d / η^d from storage. The product $\eta^c \eta^d$ is the round-trip efficiency. The loss term L is independent of dispatch and applies in every step, including idle ones.

B3 — Energy bounds

$$E^{\min} \leq E_t \leq E^{\max} \quad \forall t \in \mathcal{T} \quad (2.6)$$

where the bounds are the scenario's state-of-charge fractions scaled by nominal capacity. In the implementation these are variable bounds rather than explicit rows.

B4 — Product block coherence

Day-ahead energy settles in hourly blocks, so dispatch may not vary within a block even when the model runs at a finer resolution:

$$P_t^c = P_{t-1}^c \quad \forall t : t \bmod m \neq 0 \quad (2.7)$$

$$P_t^d = P_{t-1}^d \quad \forall t : t \bmod m \neq 0 \quad (2.8)$$

where m is the number of model steps in one product block. With a 15-minute resolution and a 60-minute day-ahead product, $m = 4$. This is also the machinery the four-hour reserve blocks of Section 3 will reuse, which is why it is expressed generally rather than hard-coded to the hourly case.

B5 — Terminal state of charge

$$E_{N-1} \geq E^{\text{term}} \quad (2.9)$$

where E^{term} is the required closing energy. The constraint exists because stored energy has no value beyond the horizon: without it the optimizer empties the battery into the final step of every window, since selling is free revenue and the consequence falls outside the model. It is applied at the far end of the window, not at the end of the committed region, so that it shapes end effects rather than the decisions actually taken.

2.6 Rolling-horizon procedure

The window model above is solved repeatedly across the price series. With $h = H/\Delta t$ steps of horizon and $c = C/\Delta t$ steps of commitment:

1. Set E_{init} to the opening state of charge.
2. Solve (2.1)–(2.9) over the next h steps.
3. Retain the first c steps of the solution as committed dispatch.
4. Set E_{init} to E_{c-1} from the solved window.
5. Advance by c steps and repeat until the series is exhausted.

The horizon exceeds the commitment ($h > c$) so that the uncommitted tail absorbs the distortion of B5. A final window shorter than one commitment step receives no terminal target, since there is no tail left to absorb it.

Setting `dispatch.mode` to `perfect foresight` collapses the procedure to a single window spanning the entire series. This reports an upper bound on achievable revenue rather than an expectation, and the same is true of the rolling mode while `dispatch.forecast.kind` remains `perfect`: each window sees its own prices exactly.

2.7 Post-optimization network charges

Network charges that fall outside the exemption of §118(6) EnWG are applied in `./src/bess_stack/finance/gr` to an already-solved dispatch, rather than by re-solving each project year under that year's tariff.

For a year y the charge is

$$\Phi_y = (\gamma_y - \gamma^{\text{priced}}) \sum_t P_t^c \Delta t + \mathbb{I}[y \text{ not exempt}] \cdot \rho \cdot \bar{P}^{\text{conn}} \cdot 1000 \quad (2.10)$$

where γ_y is the per-MWh charging fee applicable in year y , γ^{priced} is the fee the optimizer already carried in (2.1) via g^c , ρ is the successor capacity charge in EUR/kW/year and $\bar{P}^{\text{conn}}$ is the connection capacity in MW. Subtracting the fee already priced is what prevents a charge from being billed twice.

This is an approximation and its error is measured rather than assumed. Because the retained dispatch was optimized without the fee, it remains feasible but no longer optimal once the fee exists, so *revenue is understated*; and because the marginal cycles the fee should have suppressed are still present, *throughput and cycle count are overstated*. Measured against a re-solved run on a full synthetic year, a 5 EUR/MWh charging fee understates revenue by 1.0% while overstating throughput by 13.3%; at 25 EUR/MWh the figures are 50.7% and 101.1%. The asymmetry matters, because the augmentation trigger keys off cycle count and therefore fires early long before the revenue figure looks wrong. The error is exactly zero while the exemption holds.

3 Reserve Products (Specified, Not Implemented)

Nothing in this section is implemented. The parameters below are present in `scenarios/reference.yaml`, but `config.py` defines no FCR or aFRR dataclass and does not read them, and `solve_window` builds no reserve variables. What follows is the formulation these parameters imply, recorded so the schema and the intended model can be checked against each other before either is built.

3.1 Additional sets and parameters

Let $k \in \mathcal{K}$ index four-hour reserve blocks, with $\mathcal{T}_k$ the model steps in block k and H^{blk} the block length in hours.

Symbol	Description	Source field
c^{FCR}	FCR capacity price (EUR/MW/h)	<code>fcr.capacity_price_eur_per_mw_h</code>
h^{FCR}	Energy headroom held per MW awarded (h)	<code>fcr.reserve_energy_hours</code>
R^{min}	Minimum bid size (MW)	<code>fcr.min_bid_mw, afrr.min_bid_mw</code>
$c^{a\pm}$	aFRR capacity price, each direction	<code>afrr.{positive,negative}.capacity_price_eur_per_mw_h</code>
$\pi^{a\pm}$	aFRR energy price, each direction	<code>afrr.{positive,negative}.energy_price_eur_per_mwh</code>
$\alpha^\pm$	Expected activation rate	<code>afrr.{positive,negative}.activation_rate</code>
h^a	aFRR energy headroom per MW awarded (h)	<code>afrr.reserve_energy_hours</code>

3.2 Additional variables

$R_k^{\text{FCR}} \geq 0$ is the symmetric FCR award in block k ; $R_k^{a+}, R_k^{a-} \geq 0$ are the directional aFRR awards. Binary indicators $y_k^\bullet \in \{0, 1\}$ carry the minimum bid size.

3.3 Revenue terms

The objective (2.1) gains capacity revenue and expected activation revenue:

$$+ \sum_{k \in \mathcal{K}} H^{\text{blk}} \left[c^{\text{FCR}} R_k^{\text{FCR}} + c^{a+} R_k^{a+} + c^{a-} R_k^{a-} + \alpha^+ \pi^{a+} R_k^{a+} - \alpha^- \pi^{a-} R_k^{a-} \right] \quad (3.1)$$

where the sign on the negative-direction energy term reflects that π^{a-} is quoted as payment *to* the provider for absorbing energy, so a negative price is revenue. This convention must be pinned down against settlement rules before implementation; it is stated here precisely because it is the kind of sign that is easy to get backwards and hard to detect afterwards.

3.4 R1 — Minimum bid size

$$R^{\min} y_k^\bullet \leq R_k^\bullet \leq \bar{P} y_k^\bullet \quad \forall k \in \mathcal{K} \quad (3.2)$$

where $\bullet$ ranges over the three products. An award is either zero or at least $R^{\min}$, which is a semi-continuous variable and the reason reserve participation introduces binaries of its own.

3.5 R2 — Power headroom

Capacity sold must be deliverable on top of energy dispatch:

$$P_t^d + R_k^{\text{FCR}} + R_k^{a+} \leq \bar{P} \quad \forall k, \forall t \in \mathcal{T}_k \quad (3.3)$$

$$P_t^c + R_k^{\text{FCR}} + R_k^{a-} \leq \bar{P} \quad \forall k, \forall t \in \mathcal{T}_k \quad (3.4)$$

3.6 R3 — Energy headroom

Capacity must also be sustainable for the prequalification duration, which is what h^{FCR} and h^a encode:

$$E_t - (h^{\text{FCR}} R_k^{\text{FCR}} + h^a R_k^{a+}) \geq E^{\min} \quad \forall k, \forall t \in \mathcal{T}_k \quad (3.5)$$

$$E_t + (h^{\text{FCR}} R_k^{\text{FCR}} + h^a R_k^{a-}) \leq E^{\max} \quad \forall k, \forall t \in \mathcal{T}_k \quad (3.6)$$

where $h^{\text{FCR}} = 0.5$ reflects the German requirement that full FCR activation be sustainable for 30 minutes in both directions. Understating it permits the same MW to be sold to reserve and to arbitrage at once.

3.7 Known gap: activation does not move the state of charge

In the formulation above, aFRR energy enters revenue through the expected activation rate $\alpha^\pm$ but does not appear in the state-of-charge dynamics (2.4). Real activation moves energy and therefore constrains subsequent dispatch. Closing that gap requires either activation as a deterministic energy flow inside B2, or a stochastic formulation over activation scenarios. Until it is closed, reserve participation will be overvalued, because the model collects activation revenue without paying its state-of-charge consequence.